

CS244, Assignment 2  
TANGRA:  
Tangra is yet ANoother Great Recursive Acronym

Hristo Stoyanov  
stoyanov@stanford.edu

Petar Penkov  
ppenkov@stanford.edu

April 27, 2017

# Introduction

This paper aims to explore several congestion control algorithms for networks, characterized by highly varying bandwidth. The concrete use case is a two minute trace over an emulated Verizon LTE connection. All traffic is UDP, and there are no retransmissions in our setup. The setup can be thought of as a voice communication. The set of algorithms is a combination of loss and delay-based approaches that all aim for simplicity to avoid overfitting to the test trace. They include conventional algorithms like AIMD, and TCP Fast, as well as several solutions we could not match to a particular existing protocol. We measured 95-th percentile throughput and signal delay and tried to optimize for their ratio. The challenges we face include adaptation to highly varying bandwidth while maintaining stable behavior in steady state.

Our code is available at [https://github.com/rene96/cs244\\_lab2](https://github.com/rene96/cs244_lab2).

The branch master contains the Tangra code as well as the predecessor RTT-based schemes. The branch aimd contains a history of our AIMD-based schemes.

## Part A

The first step of our experiment was to compare the performance of fixed-window congestion control algorithms while ranging the window size. The values explored for the congestion window size were the default value of 50 packets and all powers of 2 from 1 to 64 packets, inclusive. We fixed the timeout to the default value of 1 second. We found that all values give consistent results with small noise through multiple runs. Delay and throughput were low for small windows and big for large windows. Furthermore, the power score improved while increasing the window size to 16, up to a maximum of 12.21, and it then started to gradually decline. The results of our experiment are summarized on Figure 1.

## Part B

To capture the variability of link bandwidth, we implemented a simple AIMD scheme. Initially we left the timeout fixed to a hardcoded value of 200 milliseconds, but later changed it to be a function of an estimated round-trip time via a low pass filter, where  $R$  is the current estimate and  $M$  is the new RTT measurement on every ACK. [1]

$$RTT \leftarrow \alpha RTT + (1 - \alpha) M$$

From this estimate of the current RTT we compute the timeout interval to  $R$ . We used the suggested parameters  $\alpha = 0.9$  and  $\beta = 2$ . While the performance of the scheme with a fixed timeout value was consistent, there was a significant bufferbloat in the latter scheme and it was deemed inadequate. The performance of the fixed-timeout scheme can be seen on Figure 2 in green. While the performance of this scheme was consistent, it failed to keep the queue mostly empty and introduced significant delay.




Figure 1: A graph of how throughput and delay change as we increase a fixed congestion window from 1 to 64 at every power of 2.

## Part C

Our first implementation of a delay-triggered congestion control scheme had two fixed threshold values. The implementation used a weighted moving average to estimate the current round-trip-time. If the RTT was below 80ms, the congestion window is increased, and if the RTT estimate is above 100ms, the window is decreased.

We experimented with several values of the thresholds, as well as different values for the increase and decrease of the windows.

## Part C1: changing the thresholds

We attempted to modify the thresholds to different values. Figure 1 shows some of the results we achieved through picking different values for the low and high water marks. These include picking 60ms-100ms, 60ms-80ms and 80ms-100ms as thresholds. This presents a very clear trade-off between delay and throughput.

## Part C2: changing the increase/decrease functions

Further, we tried to change the increase and decrease functions at each crossing of the thresholds. The best results seemed to be by achieved by using additive increase and additive decrease. Our increase function was  $2/cwnd$ . We tried to use  $5/cwnd$  as well, but that did not produce a significant improvement.




Figure 2: A graph of delay and throughput of different mechanisms.

## Part D

### Part D1: AIMDSoph

For our use case we found the simple AIMD scheme lacking in two ways. First, the timeout value of 200ms was hardcoded and would be unreasonable in situations when just the propagation delay is longer than that. In a more sophisticated version, we set the timeout to a multiple of the minimum measured round-trip time  $RT_{min}$ . We found that a delay of  $2 * RT_{min}$  works well as it implies queuing delay of one minimum RTT which is enough to

give us feedback while preventing bufferbloat. Second, it does not account for very small queuing delay. In this case the sender can probably send more packets. To address this we keep a running estimate of the RTT via a low pass filter. [1]

$$RTT \leftarrow \alpha RTT + (1 - \alpha) M$$

Here  $RTT$  denotes the current estimate and  $M$  is the measurement on the new ACK. We use the suggested parameter  $\alpha = 0.9$  to have conservative changes to the current estimate. For every ACK that maintains the running estimate drops below  $1.5 * RT_{min}$ , we do another step of additive increase, effectively doubling the additive increase gain temporarily. Furthermore, to keep the algorithm more optimistic, we decrease the multiplicative decrease factor from 2 to 1.5 ( multiply by 0.75 instead of 0.5 ) and double the additive increase gain. The results of this scheme were satisfactory, showed both low delay and decent utilization of the bottleneck, and can be seen in Figure 2 in magenta.

## Part D2: TCP Fast

TCP Fast is a delay-based algorithm that periodically updates the window size based on the average RTT [2].

$$cwnd \leftarrow (1 - \gamma) * cwnd + \gamma \left( \frac{cwnd * RT_{min}}{RTT} + \alpha \right)$$

The RTT is estimated with the low pass filter function as priorly described, and the  $RTT_{min}$  is a moving minimum. Additionally, our implementation does not limit window growth to twice the window size to allow for more aggressive adaptation to network changes. Unfortunately, we found that this algorithm cannot adapt fast enough to increasing bandwidth and results in low network utilization. The results can be seen in Figure 2 in yellow.

## Part D3: TANGRA

We further developed the idea of RTT-threshold based scheme by incorporating two more ideas. We first make the assumption that the minimum propagation time does not change over time. Given this assumption, our algorithm tracks the minimum RTT and assumes this is the propagation time. Further, we assume that the delay above the propagation time is proportional to the queueing delay. We write a control-loop that adjusts the congestion window based on these two assumptions.

The second idea we had was to adjust the congestion window in a way that is proportional to the difference between the current RTT estimate and a target RTT. This decision was motivated by the idea that the further away from the target the current state of the system is, the bigger the correction we need. We define our target as

$$RTT_{target} \leftarrow t * RTT_{min}$$

where  $t$  is an adjustable parameter which specifies how big the queueing delay should be relative to the propagation delay. Our best results are achieved with  $t = 1.4$ . In addition,

we define  $t_{bound}$  and  $t_{high}$ , as additional thresholds which will serve to adjust the congestion window. In our default implementation we used  $t_{bound} = 1.6$  and  $t_{high} = 1.8$ .

This results in 4 levels where the RTT estimate may fall:

1.  $RTT > t_{high} * RTT_{min}$ . The current congestion window is overshooting the target and we need to reduce it drastically.
2.  $t_{high} * RTT_{min} \geq RTT > t_{bound} * RTT_{min}$ . This serves as an interval where a smaller reduction to the congestion window is applied.
3.  $t_{bound} * RTT_{min} \geq RTT > t * RTT_{min}$ . If this condition is true, the congestion window is not changed.
4.  $t * RTT_{min} > RTT$  implies that the current congestion window is too small, and hence, we increase it.

In order to estimate the  $RTT$ , we again use a low pass filter, with  $\alpha = 0.7$ . This may seem too aggressive, but we opted to use that after running experiments with different values as it provided a trade-off between performance on the microbenchmark and stability in a network with variance that in RTT over time. Figure 3 shows several runs with different low pass filters.

The best score our implementation achieved was 40.74 at a throughput of 3.83 Mbps and 95<sup>th</sup> percentile signal delay of 94 ms. Although this was the best score we achieved over multiple runs of our implementation, this score is hardly reproducible. Most runs of the benchmark result in a scores with throughput in the range of 3.75 – 3.85 Mbps and delays of 94 – 105 ms.

Although our implementation performs very well on this tiny benchmark, this approach to congestion control has several drawbacks. First, it is very unclear how several flows will interact with each other. A scenario worth exploring is where two or more flows share a bottleneck, yet their respective paths are different and the corresponding proportional queues the flows try to maintain will be different. This may lead to a much different allocation of the bottleneck than would otherwise occur with regular TCP.

Another drawback of this approach is the fact that if the  $RTT_{target}$  is set too high relative to the available buffer at the bottleneck link, the algorithm will likely produce a high loss rate and waste traffic at the links before the bottleneck by constantly attempting to increase its congestion window.

Future work may be able to address these problems by creating a mechanism to dynamically adjust the target, while using the same idea of proportionality to try to hit that target.

## Part E: Naming TANGRA

Tangra stands for “Tangra is yet ANother Great Recursive Acronym.” It also has a nice ring to it.




Figure 3: Multiple runs of the simulation at  $t = 1.4$  and different low pass  $\alpha$  parameters.

## References

- [1] V. Jacobson. “Congestion Avoidance and Control”. In: *SIGCOMM Comput. Commun. Rev.* 18.4 (Aug. 1988), pp. 314–329. ISSN: 0146-4833. DOI: 10.1145/52325.52356. URL: <http://doi.acm.org/10.1145/52325.52356>.
- [2] D. Wei. “FAST TCP: Motivation, Architecture, Algorithms, Performance”. In: *IEEE/ACM Transactions on Networking* 14.6 (Dec. 2006), pp. 1246–1259. ISSN: 1558-2566. DOI: 10.1109/TNET.2006.886335. URL: <http://ieeexplore.ieee.org/document/4032738/>.